




RPC-Bench: A Fine-grained Benchmark for Research Paper Comprehension

Yelin Chen, Fanjin Zhang, Suping Sun, Yunhe Pang, Yuanchun Wang, Jian Song, Xiaoyan Li, Lei Hou, Shu Zhao, Jie Tang, Juanzi Li

Xinjiang University, Renmin University of China, Anhui University, Sun Yat-sen University,

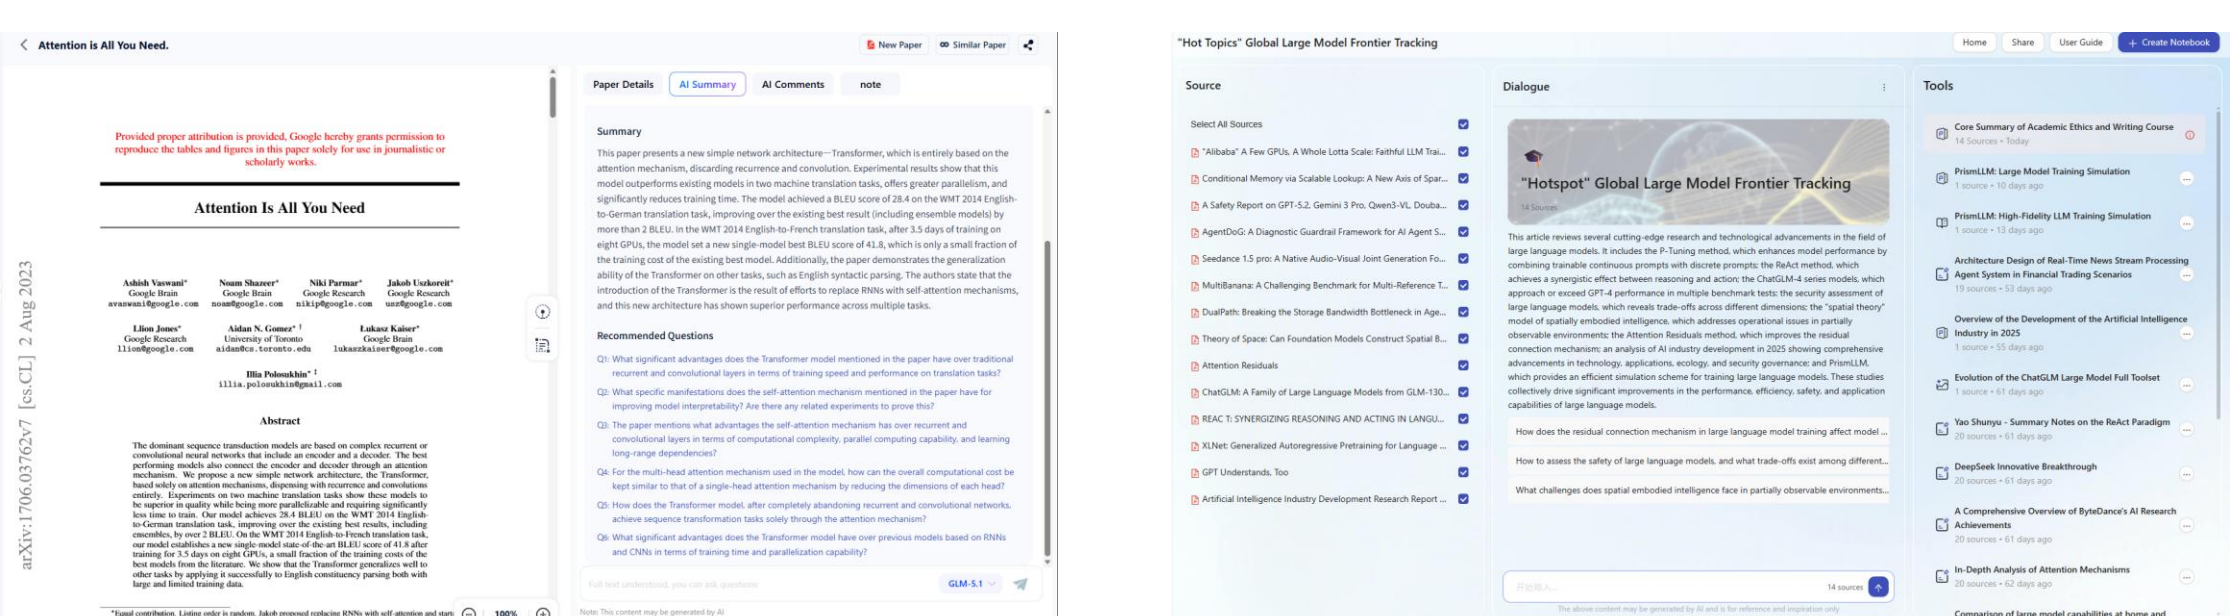
Z.ai, Tsinghua University, University of Southampton

Project Page: <https://rpc-bench.github.io/>



Why Research Paper Comprehension?

Research paper comprehension is a key prerequisite for **reliable research copilots** and **subsequent scientific discovery**.



AI Paper Reading

Research Labs

Related Benchmarks

- Limited scale and paper coverage
- Rely on synthetic QA pairs
- Lack fine-grained comprehension taxonomy
- Provide narrow evaluation metrics
- Do not jointly support text and visual inputs

Benchmarks	Papers	QA	Real QA	Taxonomy	Eval. Metrics	Textual inp.	Visual inp.
PeerQA	208	579	✓	task	Corr.	✓	✗
SPIQA	25.5K	270K	✓	task	LLMLogScore	✓	✓
ArXivQA	16.6K	-	✗	task	-	✗	✓
DocGenome	500K	-	✗	task	GPT-acc	✗	✓
RPC-Bench	4150	61.3K	✓	content	Conc., F1-like	✓	✓

Table 1: Comparison with relevant research paper Benchmarks. Conc.=Conciseness; Corr.=Correctness; F1-like is defined as the harmonic mean of correctness and completeness; inp.=input. "Eval. Metrics" are LLM-based metrics.

Evaluation Protocol

- Uses **accuracy** for binary claim verification and **LLM-as-a-Judge** scoring for open-ended QA.
- Scores correctness, completeness, and conciseness on a 0-5 scale; derives F1-like from correctness-completeness and informativeness by penalizing verbose answers.

$$\text{F1-like} = \frac{(1 + \beta^2) \times (\text{Correctness} \times \text{Completeness})}{\beta^2 \times \text{Correctness} + \text{Completeness}}$$

$$\text{Informativeness} = \text{F1-like} \times \frac{\text{Conciseness}}{5}$$

Adopts the most human-aligned setup: **title/abstract-augmented prompts**, **separate dimension scoring**, **decimal scores**, and **averaged GPT-5/Gemini-3-pro judgments**, selected by Bradley-Terry correlation and pairwise AUC


Setting	P-BT	S-BT	PW-AUC	Avg.
AUG+SEP	0.8955	0.9137	0.7125	0.8406
AUG+JOI	0.9003	0.9091	0.6966	0.8353
RAW+SEP	0.8773	0.9137	0.7054	0.8321
RAW+JOI	0.8759	0.9137	0.7100	0.8332

Table 4: Agreement between human and model judgment w.r.t. prompt configurations. (AUG: enhancement with title and abstract; SEP/JOI: separate/joint evaluation)

Model	P-BT	S-BT	PW-AUC	Avg.
GPT-5 _{int}	0.9059	0.9227	0.7027	0.8437
Claude-4.5 _{int}	0.8928	0.9136	0.7038	0.8367
Gemini-3 _{int}	0.9091	0.9227	0.7090	0.8469
GPT-5 _{dec}	0.9213	0.9137	0.7255	0.8535
Claude-4.5 _{dec}	0.8873	0.9091	0.7197	0.8387
Gemini-3 _{dec}	0.9170	0.9182	0.7330	0.8561

Table 5: Agreement between different LLM judges and human assessments. Subscript *int* denotes integer scoring, and subscript *dec* denotes decimal scoring.

RPC-Bench: Construction Pipeline



- Built from **real review-rebuttal interactions** through **decomposition, rewriting, filtering, and human verification**.
- Each paper is provided in **structured-text** and **rendered-page** formats for LLM/VLM evaluation.

Taxonomy Design

- Designed around the natural **research flow** of academic papers, progressing from **what** and **how** to **why** questions to trace the full logic of a paper.
- Includes **free-form QA** on concepts, methods, and experiments, plus **binary claim verification** for judging whether paper-grounded claims are true or false.

Task Taxonomy

- Concept Understanding (C.U.) [What-4.27%]: Clarifies or explains key concepts, terminology, theoretical viewpoints, or information conveyed in figures, tables, or formulas.
- Methods
 - Method Disambiguation (M.D.) [What-8.91%]: Clarifies methodological details to resolve misunderstandings or ambiguities, ensuring an accurate understanding of proposed approaches.
 - Method Mechanics (M.M.) [How-9.91%]: Questions about the implementation or function of methodological workflow or components, such as the effect of specific modules in models.
 - Motivation Analysis (M.A.) [Why-6.29%]: Explores the rationale behind a proposed method or decision.
 - Method Comparison (M.C.) [11.75%]: Compares the proposed approach with baseline methods, analyzing similarities, differences, or performance to highlight novelty.
- Experiments
 - Experimental Exposition (E.E.) [What-13.07%]: Explains results, infers how changes to experiments could affect results or conclusions, and handling reasoning tasks (i.e., calculations, counting, or comparisons).
 - Experimental Setup (E.S.) [How-6.77%]: About the design, configuration, and execution of experiments.
 - Experimental Analysis (E.A.) [Why-14.08%]: Studies the reasons of specific experimental results, links them to the proposed approach, and assesses their generalizability and potential impact.
- Claim Verification (C.V.) [24.95%]: Binary classification tasks that judge whether claims, hypotheses, or experimental conclusions are correct.

RPC-Bench Overview

Statistics	train	val	test
Papers	3100	850	200
Accept	1980	609	116
year	2013-2021	2022-2023	2024
Venue	12	7	4
QA	45651	12895	2787
A/M Q	25.2/105	26.2/261	24.2/250
A/M A	70.5/297	121.5/1337	87.9/773

Table 2: Statistics of the RPC-Bench. A/M Q: average/max question length. A/M A: average/max answer length. Lengths are measured in words.




Figure 2: Domain distribution of RPC-Bench. ML: Machine Learning; CV: Computer Vision; NLP: Natural Language Processing; RL: Reinforcement Learning.

Experimental Results

- Traditional metrics fail to reflect true paper comprehension; correctness, completeness, and conciseness reveal clear gaps even for top models.
- Text-only LLMs outperform image-based VLMs, while small DCM/RAG models struggle with long-context reasoning.
- Models handle concept/method questions better than experimental analysis and claim verification, which require evidence integration and strict factual judgment.

Model Type	Model	Traditional		LLM-as-a-judge				
		R-L	BERTS.	Concise.	Correct.	Complete.	F1-like	Info.
LLM	DeepSeek-V3.2	19.22	<u>55.60</u>	56.31	58.73	55.19	56.91	32.04
	GLM-4.7	17.09	48.58	54.34	54.36	51.75	53.02	28.81
	Qwen3	16.16	54.25	41.44	55.88	56.64	56.26	23.31
	GPT-5	16.89	54.52	54.93	69.10	67.33	68.20	37.46
	Claude-4	16.60	54.02	41.37	58.53	58.44	58.48	24.19
	Gemini-2.5	18.24	55.67	54.87	<u>62.65</u>	<u>59.03</u>	<u>60.79</u>	33.35
DCM	DocOwl2(V)	14.32	46.42	50.19	11.75	6.66	8.50	4.27
	Docopilot(V)	16.92	53.82	39.31	18.31	17.12	17.69	6.96
	Monkey(V)	20.16	55.19	54.61	17.08	11.27	13.58	7.41
VLM	GLM-4.6V	<u>19.38</u>	54.76	64.55	47.32	43.43	45.29	29.23
	Qwen3(V)	14.70	53.72	22.64	20.17	20.14	20.16	4.56
	GPT-5(V)	17.32	54.85	61.47	58.90	55.34	57.07	<u>35.08</u>
	Claude-4(V)	13.33	50.63	31.63	54.16	53.32	53.74	16.99
	Gemini-2.5(V)	17.27	54.85	51.71	48.39	45.59	46.95	24.28
RAG	HippoRAG2	18.71	54.16	45.77	33.13	27.88	30.28	13.86
	MemoRAG	13.55	52.70	51.31	24.19	19.10	21.35	10.96
	Raptor	18.35	54.00	36.47	25.28	20.82	22.84	8.33
	VdocRAG(V)	17.77	52.22	<u>61.54</u>	21.17	13.88	16.77	10.32
	VisRAG(V)	16.80	54.93	39.90	26.24	23.63	24.87	9.92

Table 3: Evaluation results of free-form QA on the test set. R-L=ROUGE-L; BERTS.=BERTScore; Concise.=Conciseness; Correct.=Correctness; Complete. = Completeness; Info. = Informativeness. The best results are highlighted in **bold**, and the second-best results are underlined.




Figure 4: Comparison of LLMs and VLMs on open-ended question answering (F1-like score; left), and the performance of all models on claim verification tasks (ACC; right).

Conclusion

- Introduce RPC-Bench, a large-scale benchmark with **4,150** papers, **61.3K** QA pairs, and **9** fine-grained categories aligned with the research workflow.
- Transform real review-rebuttal exchanges into high-quality QA pairs through an LLM-human collaborative annotation framework.
- Provide a comprehensive evaluation protocol that jointly measures correctness, completeness, and conciseness, with strong alignment to human judgments.
- Benchmark **28** advanced models across LLMs, VLMs, DCMs, and RAG systems in both text and page-image settings.